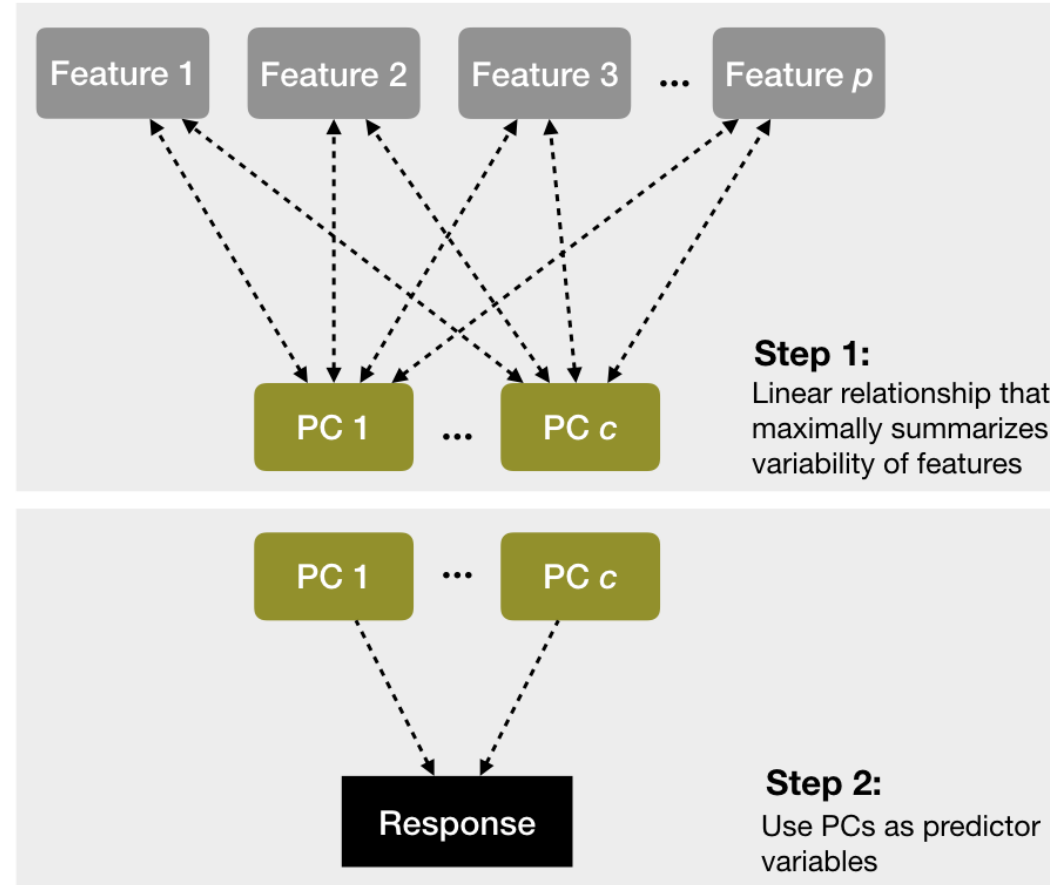


Supervised Learning

Principal component regression and partial least squares

July 6th, 2021

Principal component regression (PCR)



Example data: NFL teams summary

Created dataset using `nflfastR` summarizing NFL team performances from 1999 to 2020

```
library(tidyverse)
nfl_teams_data <- read_csv("http://www.stat.cmu.edu/cmsac/sure/2021/materials/data/regression_pro
nfl_model_data <- nfl_teams_data %>%
  mutate(score_diff = points_scored - points_allowed) %>%
  # Only use rows with air yards
  filter(season >= 2006) %>%
  dplyr::select(-wins, -losses, -ties, -points_scored, -points_allowed, -season, -team)
nfl_model_data
```

```
## # A tibble: 480 × 49
##   offense_com...1 offen...2 offen...3 offen...4 offen...5 offen...6 offen...7 offen...8 offen...9
##           <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1         0.561     3662     1350      6.40      3.28     4284      8.01     1582      4.94
## 2         0.480     2371     2946      5.10      5.56     4698     11.3       942      4.22
## 3         0.612     3435     1667      6.41      3.74     4082      7.88     1391      4.24
## 4         0.564     2718     1555      5.70      3.73     3833      8.91     1243      4.62
## 5         0.569     3264     1674      5.72      4.10     4348      8.07     1553      4.79
## 6         0.525     3286     1940      6.12      4.02     4564      8.90     1374      4.89
## 7         0.588     3827     1648      6.88      3.91     5064      9.76     1466      4.48
## 8         0.565     2893     1347      5.16      3.69     3766      7.43     1533      4.84
## 9         0.569     3838     1954      7.04      4.23     4681      9.38     1427      4.63
```

Implement PCR with **pls** package

Similar syntax to `lm` formula but specify the number of PCs (`ncomp`)

```
library(pls)
nfl_pcr_fit <- pcr(score_diff ~ ., ncomp = 2, scale = TRUE, data = nfl_model_data)
summary(nfl_pcr_fit)
```


```
## Data:      X dimension: 480 48
##      Y dimension: 480 1
## Fit method: svdpc
## Number of components considered: 2
## TRAINING: % variance explained
##           1 comps  2 comps
## X           21.49   41.62
## score_diff   84.01   87.49
```

Tuning PCR with **caret**

To perform PCR we need to tune the number of principal components

- Tune # components in PCR with **caret**
- train with 10-fold CV using pcr from **pls**

```
set.seed(2013)
library(caret)
cv_model_pcr <- train(
  score_diff ~ .,
  data = nfl_model_data,
  method = "pcr",
  trControl = trainControl(method = "cv", num
  preProcess = c("center", "scale"),
  tuneLength = ncol(nfl_model_data) - 1)
ggplot(cv_model_pcr) + theme_bw()
```



Tuning PCR with **caret**

By default returns model with minimum CV error as `finalModel`

```
summary(cv_model_pcr$finalModel)
```

```
## Data:      X dimension: 480 48
##      Y dimension: 480 1
## Fit method: svdpc
## Number of components considered: 43
## TRAINING: % variance explained
##           1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X           21.49   41.62   53.04   61.70   66.63   70.79   74.5
## .outcome    84.01   87.49   87.75   89.69   89.69   90.06   90.2
##           8 comps  9 comps 10 comps 11 comps 12 comps 13 comps 14 comps
## X           77.81   80.56   83.10   85.27   87.18   88.99   90.43
## .outcome    90.31   90.32   90.66   90.82   91.03   91.05   91.12
##          15 comps 16 comps 17 comps 18 comps 19 comps 20 comps 21 comps
## X           91.85   93.10   94.00   94.76   95.49   96.18   96.79
## .outcome    91.15   91.28   91.28   91.28   91.28   91.34   91.71
##          22 comps 23 comps 24 comps 25 comps 26 comps 27 comps 28 comps
## X           97.30   97.77   98.21   98.57   98.84   99.07   99.29
## .outcome    91.73   91.80   91.95   91.96   92.64   92.72   92.79
##          29 comps 30 comps 31 comps 32 comps 33 comps 34 comps 35 comps
## X           99.47   99.61   99.74   99.85   99.92   99.98   99.99
## .outcome    92.81   92.81   92.81   92.81   92.81   92.81   92.81
```

Tuning PCR with **caret**

Modify selectionFunction in train to be the oneSE rule

```
set.seed(2013)
cv_model_pcr_onese <- train(
  score_diff ~ .,
  data = nfl_model_data,
  method = "pcr",
  trControl =
    trainControl(method = "cv", number = 10,
                  selectionFunction = "oneSE")
  preProcess = c("center", "scale"),
  tuneLength = ncol(nfl_model_data) - 1)
```

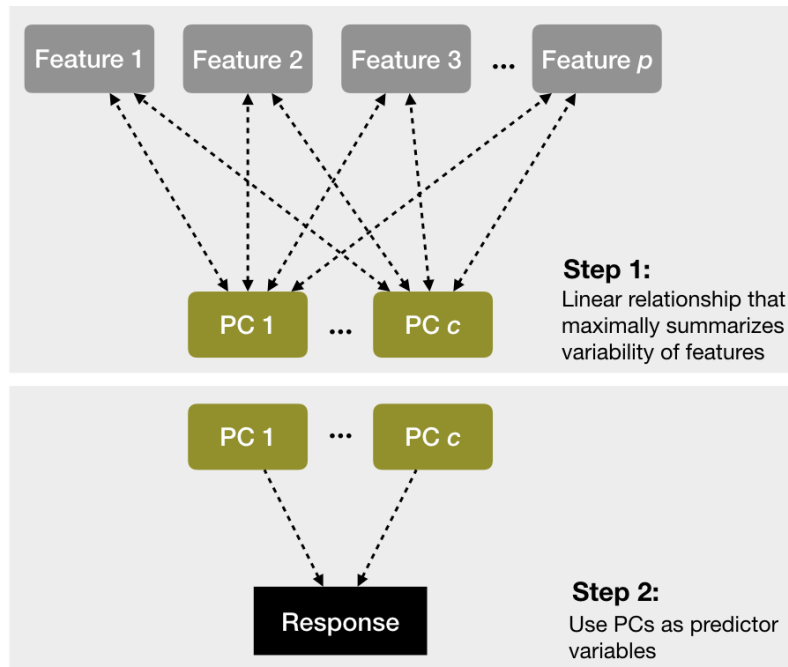
```
summary(cv_model_pcr_onese$finalModel)
```

```
## Data:      X dimension: 480 48
##      Y dimension: 480 1
## Fit method: svdpc
## Number of components considered: 32
## TRAINING: % variance explained
##           1 comps  2 comps  3 comps  4 comps  5 co
## X           21.49   41.62   53.04   61.70   66
## .outcome    84.01   87.49   87.75   89.69   89
##           8 comps  9 comps 10 comps 11 comps 12
## X           77.81   80.56   83.10   85.27
## .outcome    90.31   90.32   90.66   90.82
##           15 comps 16 comps 17 comps 18 comps
## X           91.85   93.10   94.00   94.76
## .outcome    91.15   91.28   91.28   91.28
##           22 comps 23 comps 24 comps 25 comps
## X           97.30   97.77   98.21   98.57
## .outcome    91.73   91.80   91.95   91.96
##           28 comps 29 comps 30 comps 31 comps 32 comps
```

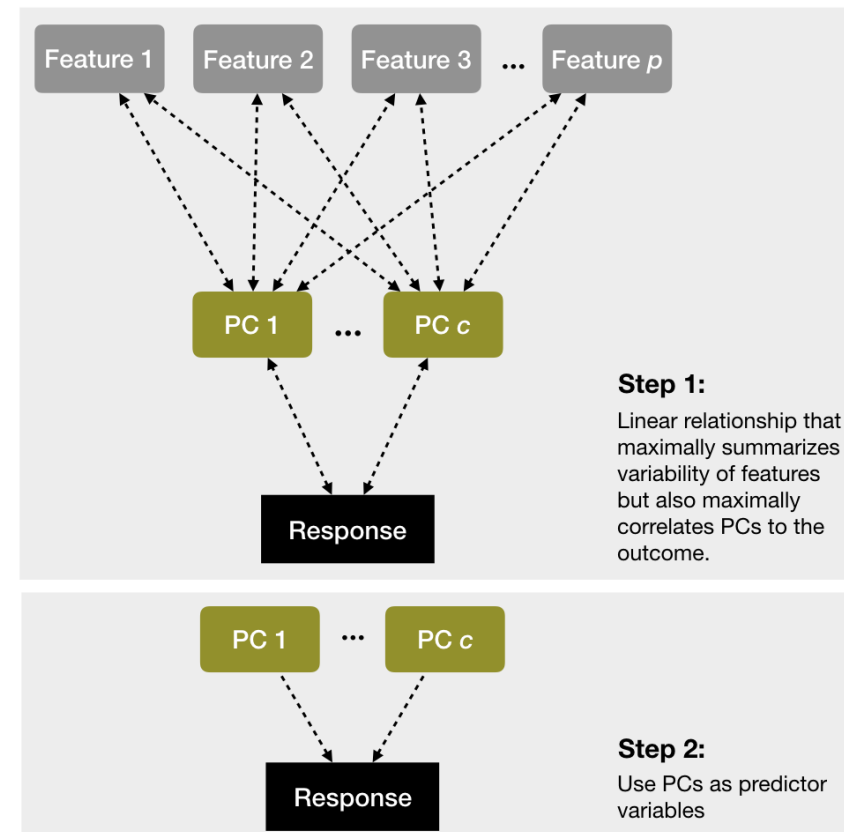
Partial least squares (PLS)

PCR is agnostic of response variable

(a) Principal Components Regression



(b) Partial Least Squares Regression



PLS as supervised dimension reduction

First principal component in PCA:

$$Z_1 = \phi_{11}X_1 + \phi_{21}X_2 + \cdots + \phi_{p1}X_p$$

In PLS we set ϕ_{j1} to the coefficient from **simple linear regression** of Y on each X_j

- Remember this slope is proportional to the correlation! $\hat{\beta} = r_{X,Y} \cdot \frac{s_Y}{s_X}$
- Thus Z_1 in PLS places most weight on variables strongly related to response Y

To compute Z_2 for PLS:

- Regress each X_j on Z_1 , residuals capture signal not explained by Z_1
- Set ϕ_{j2} to the coefficient from **simple linear regression** of Y on these residuals for each variable

Repeat process until all Z_1, Z_2, \dots, Z_p are computed (**PLS components**)


Then regress Y on Z_1, Z_2, \dots, Z_{p^*} , where $p^* < p$ is a tuning parameter

Tuning PLS with **caret**

```
set.seed(2013)
cv_model_pls <- train(
  score_diff ~ .,
  data = nfl_model_data,
  method = "pls",
  trControl =
    trainControl(method = "cv", number = 10,
                  selectionFunction = "oneSE")
  preProcess = c("center", "scale"),
  tuneLength = ncol(nfl_model_data) - 1)
ggplot(cv_model_pls) + theme_bw()
```

Sharp contrast with PCR results!

Fewer PLS components because they are guided by the response variable



But how do we summarize variable relationships without a single coefficient?

Variable importance with **vip** package


Variable importance attempts to quantify how influential variables are in the model

- e.g., absolute value of t -statistic in regression

For PLS: weighted sums of the absolute regression coefficients across components

- Weights are function of reduction of RSS across the number of PLS components

```
# Check out `cv_model_pls$finalModel$coeffici  
library(vip)  
vip(cv_model_pls, num_features = 10,  
    method = "model") +  
theme_bw()
```



Partial dependence plots (PDP) with **pdp** package

PDPs display the change in the average predicted response as the predictor varies over their marginal distribution

- More useful for non-linear models later on!

```
library(pdp)  
partial(cv_model_pls, "offense_total_epa_pass", plot = TRUE)
```

