

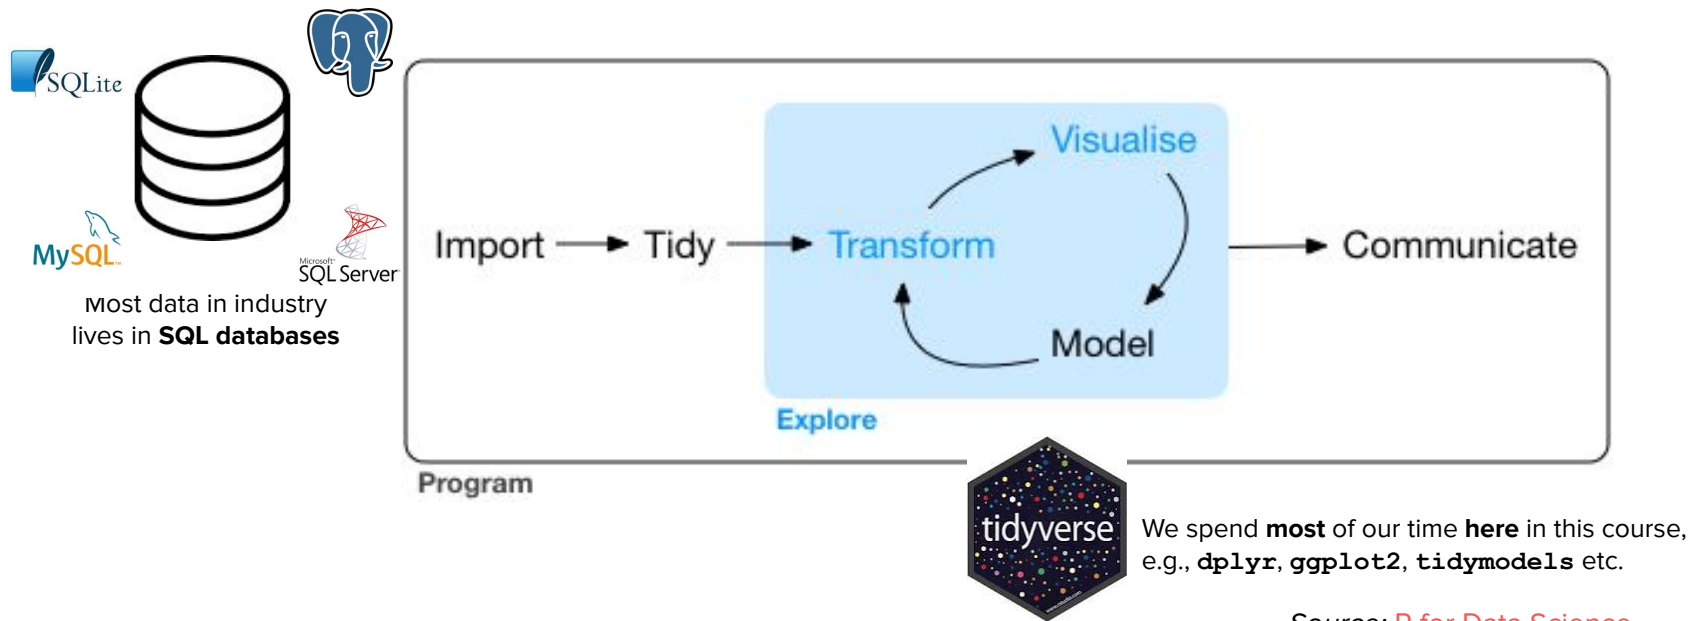
Data Engineering - Lecture 5

A practical approach to SQL - Part 1

Shamindra Shrotriya (CMU)

So what does a typical **data-driven** *workflow* look like?

Data-driven workflows adopt an interactive pipeline



Source: [R for Data Science](#)

Takeaway: being able to **efficiently extract SQL data** is vital for success

Aren't R/python/Julia **alone** sufficient for this purpose?

No - But they work brilliantly **with SQL**!

SQL databases allow you to **persistently store** and **organize** data

Support a streamlined **Extract-Transform-Load (ETL)** process for streaming data

Provide **access management restrictions** to **specific data**, e.g., health records

Allow for **explicit linkages across tables** (primary and foreign keys)

Enable **indexes** to be defined on tables for efficiency, e.g., **date/time** fields

Takeaway: use R for **accessing subsets** of data from a **SQL** database for modeling

Key idea **query**: *table(s) → table*

SQL provides a consistent grammar (**Structured Language**) for asking and answering questions (**Queries**) about your collected data

SQL tables are nouns, on which you ask targeted queries

← Columns (variables) →

↑ Observations (rows) ↓

	dest	month	day	mnd	mxl	avd
1	ABQ	12	1	-36	-36	-36
2	ABQ	12	2	-17	-17	-17
3	ABQ	12	3	20	20	20
4	ABQ	12	4	27	27	27
5	ABQ	12	5	32	32	32
6	ABQ	12	6	46	46	46
7	ABQ	12	7	53	53	53
8	ABQ	12	8	114	114	114
9	ABQ	12	9	57	57	57
10	ABQ	12	10	108	108	108

Tables are just **2D representations** of data

A **collection** of **columns** and **observations**

These are similar to data **frames/tibbles** in R

“**tibble**” even phonetically **sounds like “table”**

You’re already used to them in R - yay!

Takeaway: **data frames** in R/Python are natural analogues of **SQL** tables

SQL grammar comes pre-built with common keywords

	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time
1	2013	1	1	517	515	2	830
2	2013	1	1	533	529	4	850
3	2013	1	1	542	540	2	923
4	2013	1	1	544	545	-1	1004
5	2013	1	1	554	600	-6	812
6	2013	1	1	554	558	-4	740
7	2013	1	1	555	600	-5	913
8	2013	1	1	557	600	-3	709
9	2013	1	1	557	600	-3	838
10	2013	1	1	558	600	-2	753



Thousands **more observations**

SQL Code

```
SELECT dest, month, day,  
       MIN(arr_delay) AS mnd,  
       MAX(arr_delay) AS mxd,  
       AVG(arr_delay) AS avd  
FROM flights  
GROUP BY dest, month, day  
ORDER BY dest, month DESC,  
         day  
LIMIT 10;
```

	dest	month	day	mnd	mxd	avd
1	ABQ	12	1	-36	-36	-36
2	ABQ	12	2	-17	-17	-17
3	ABQ	12	3	20	20	20
4	ABQ	12	4	27	27	27
5	ABQ	12	5	32	32	32
6	ABQ	12	6	46	46	46
7	ABQ	12	7	53	53	53
8	ABQ	12	8	114	114	114
9	ABQ	12	9	57	57	57
10	ABQ	12	10	108	108	108

Takeaway: these keywords (verbs) allow you to systematically query tables (nouns)

SQL keywords have a direct **bidirectional** to **dplyr** verbs

SELECT	↔	<code>select()</code> , <code>mutate()</code> , <code>summarize()</code>	
FROM	↔	specified input data frame/tibble	
WHERE	↔	<code>filter()</code>	
GROUP_BY	↔	<code>group_by()</code>	
HAVING	↔	<code>group_by() %>% summarize() %>% filter()</code>	
ORDER BY	↔	<code>arrange()</code>	
LIMIT	↔	<code>"head()"</code> or <code>"tail()"</code>	Adapted from: Ian Cook

Takeaway: **dplyr** developed this precise relationship to **SQL** **by design** over time

A reminder as to why I use SQL

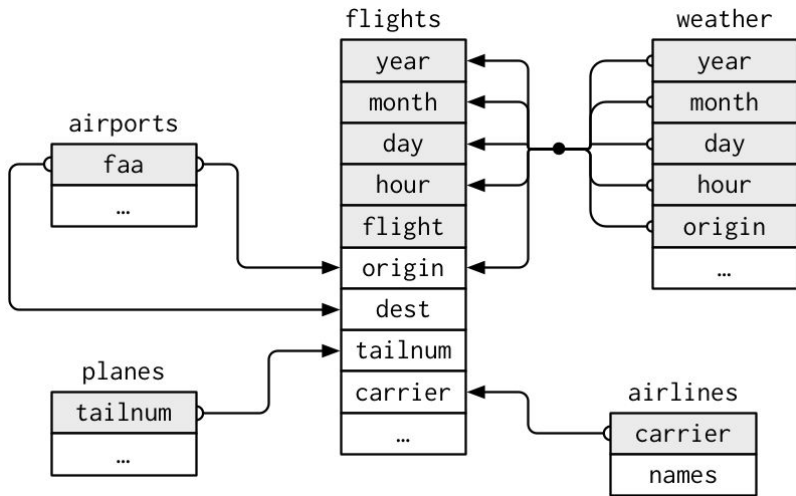
I like using **SQL** because it's *fun* and *necessary*

Specifically **SQL** allows me to **ask** and **answer** precise **questions** on collected data, in a manner that is both easy to *communicate* and *scales* with data size.

Always first aim to visualize your database before using SQL

We'll use the **nycflights13** database for our analysis

What: Contains flight info for NYC departures to various US destinations in 2013



Source: [nycflights13](#)

flights: all NYC departures in 2013

weather: hourly data for each airport

planes: construction info for each plane

airports: airport names and locations

airlines: two letter carrier codes/names

Takeaway: building this **mental picture** *up front* gets us in the right **SQL mindset**

Let's run sqlite3 queries **within R** for **nycflights13**

sqlite: “small, fast, self-contained, high-reliability, full-featured, SQL database engine”

```
> install.packages(c("dittodb", "RSQLite", "nycflights13"))  
  
> NYC_CONN <- DBI::dbConnect(RSQLite::SQLite(), ":memory:")  
  
> dittodb::nycflights13_create_sql(NYC_CONN)  
  
> fetch_query <- function(query, con = NYC_CONN) {  
  return(DBI::dbGetQuery(con, query))  
}  
  
> fetch_query("SELECT * FROM flights LIMIT 11")
```

SELECT ↔ `dplyr::select()`

We can **SELECT** any column we want from a table

Answer to: how can we select specific columns from a table

```
> SELECT <column_name> FROM <table_name>
```

Let's glimpse 10 rows and all variables from the flights data

```
> SELECT * FROM flights LIMIT 10;
```

The ***** means return all (any) columns

SQL will return any **10** rows, so the original **flights** order ***may not be*** preserved

Takeaway: don't **assume** that SQL results **implicitly preserve** original data **ordering**

We can **SELECT** any column we want from a table (cont'd)

```
> SELECT dep_time, arr_time, flight FROM flights LIMIT 10;
```

The equivalent **dplyr** code is

```
> flights %>% select(dep_time, arr_time) %>% head(10)
```

Note that **original flights** ordering *is* preserved in **dplyr**

SQL operates on **sets** of observations, which are an unordered collection

We'll later control ordering explicitly in **SQL** using **ORDER BY**

Takeaway: always add a **LIMIT** clause when you are **just** selecting from a table

SELECT ↔ `dplyr::mutate()`

We can also use **SELECT** to create new variables

Answer to: how can add new columns to a table, e.g., from existing ones?

Let's get a measure of average speed (miles per hour) for each flight

```
> SELECT flight, distance/(air_time/60) AS speed FROM flights LIMIT 10;
```

We created the required column and named it **AS speed**

In **dplyr** we have the **mutate()** verb

```
> flights %>% mutate(speed = distance/(air_time/60)) %>% select(flight,  
speed) %>% head(10)
```

Takeaway: **SELECT** serves to **pick existing** columns or to **create new** ones

SELECT ↔ `dplyr::summarize()`

We can also aggregate on columns using **SELECT**

Answer to: how can create summary statistics across **all** rows?

SQL has built in aggregate functions: **MIN**, **MAX**, **COUNT**, **SUM**, **AVG**, ...

```
> SELECT MIN(air_time) AS min_ar, MAX(air_time) AS max_ar from flights;
```

We didn't need **LIMIT** here, since we **returned a single** aggregate observation

We can get the **total number of observations** using **COUNT (*)** operator

```
> SELECT COUNT(*) AS num_obs from flights;
```

Takeaway: Aggregations are **most effective** when working across groups of data

WHERE ↔ `dplyr::filter()`

We can filter observations **WHERE** a criteria is met

Answer to: how can we subset observations which meet a given criteria?

Fetch all flights which departed from “JFK” (but limit to 10 observations)

```
> SELECT * FROM flights WHERE origin = "JFK" LIMIT 10;
```

Count flights which did not arrive at “JFK”

```
> SELECT COUNT(*) FROM flights WHERE dest != "JFK";
```

We can also use these comparison operators `=`, `!=`, `<`, `<=`, `>`, `>=`

Takeaway: Filtering operations in **SQL** are similar to **R**, except `==` is just `=` in SQL

How about **WHERE** a variable is IN or NOT IN a range?

Find 20 records which have a tail number matching either {"N593JB", "N532UA"}

```
> SELECT * FROM flights WHERE origin IN ("N593JB", "N532UA") LIMIT 20;
```

Flights which did not depart in either {Dec, Jan} and had an arrival delay > 120 mins

```
> SELECT * FROM flights WHERE month NOT IN (1, 12) AND arr_delay > 120  
LIMIT 10;
```

We could have written the following in dplyr

```
> flights %>% filter(!(month %in% c(1, 12)) & arr_delay > 120) %>% head(20)
```

Takeaway: It's helpful to re-write queries in **R**, and pattern match to **SQL**

Missing values are **NULL** in **SQL** and dealt with differently

Get weather records where wind gust is not missing

```
> SELECT * FROM weather WHERE wind_gust IS NOT NULL LIMIT 20;
```

Note: `wind_gust != NULL` does **not work**, **NULL** values don't match this way

In **R**, missing values are **NA** so we could do either of the following in **dplyr**

```
> weather %>% filter(!is.na(wind_gust))
```

```
> weather %>% drop_na(wind_gust) %>% head(20)
```

Takeaway: Be **careful** when dealing with missing (**NULL**) values in **SQL**

GROUP BY ↔ `dplyr::group_by()`

We can **GROUP BY** variables and do aggregate calculations

Answer to: how can we compute aggregate summaries by groups across columns?

Get average arrival delay by flight origin

```
> SELECT origin, AVG(arr_delay) AS avd FROM flights GROUP BY origin;
```

Note that we renamed the average arrival delay column **AS avd**

In **dplyr** we could do the following

```
> flights %>% group_by(origin) %>% summarize(avd = mean(arr_delay, na.rm = TRUE))
```

Takeaway: similar verbs have slightly different implementations in **R** and **SQL**

We can also **GROUP BY** multiple variables

Get minimum, maximum, and average arrival delay by month day and destination

```
> SELECT dest, month, day,  
        MIN(arr_delay) AS mnd,  
        MAX(arr_delay) AS mxd,  
        AVG(arr_delay) AS avd  
FROM flights  
GROUP BY dest, month, day  
LIMIT 10;
```

Takeaway: SQL handles the variable groups, you specify **which** variables to group

HAVING ↔ `dplyr::group_by()` %>%

`dplyr::summarize()` %>%

`dplyr::filter()`

We can filter aggregated values **HAVING** met a condition

Answer to: how can filter on the aggregated values?

Given number of plane engines, how many had more less than 200 manufacturers?

```
> SELECT engines, COUNT(*) AS tot_num  
  
FROM planes  
  
GROUP BY engines  
  
HAVING tot_num < 200;
```

We could have done **HAVING COUNT(*) < 200;**

We can filter aggregated values **HAVING** met a condition

Given number of plane engines, how many had more less than 200 manufacturers?

In **dplyr** we could do

```
> planes %>% group_by(engines) %>%  
  summarize(tot_num = n()) %>% filter(tot_num < 200)
```

Or we could use the nice **count** verb to avoid an explicit **group_by/filter**

```
> planes %>% count(engines, name = "tot_num") %>% filter(tot_num < 200)
```

ORDER BY ↔ `dplyr::arrange()`

We can **ORDER BY** many columns for displaying output

Answer to: how to **display** tables **sorted** by one or more columns?

Get minimum, maximum, and average arrival delay by month day and destination

```
> SELECT dest, month, day,  
        MIN(arr_delay) AS mnd, MAX(arr_delay) AS mxd,  
        AVG(arr_delay) AS avd  
FROM flights  
GROUP BY dest, month, day  
ORDER BY dest, month DESC, day  
LIMIT 10;
```

Takeaway: ordering is by **default ascending**, unless you specify descending

So ***what's*** next...?

So much more - but we'll aim for the following

Table aliases: shorthand ways to reference specific tables in your queries

Subqueries: queries within queries

JOINS: how to connect information across tables

WINDOW functions: how to run non-aggregated operations across groups

References

Wickham, Hadley, Mine Çetinkaya-Rundel, and Garrett Grolemund. *R for data science*. " O'Reilly Media, Inc.", 2023. [[Link](#)]

Wickham H (2022). *nycflights13: Flights that Departed NYC in 2013*. R package version 1.0.2, [[Link](#)]

Cook, Ian. *tidyquery and queryparser: Translating SQL Queries to dplyr Pipelines* [[Link](#)]

Teate, Renee MP (2021). *SQL for data scientists: a beginner's guide for building datasets for analysis*. [[Link](#)]

Evans, Julia *Become a SELECT star* [[Link](#)]