

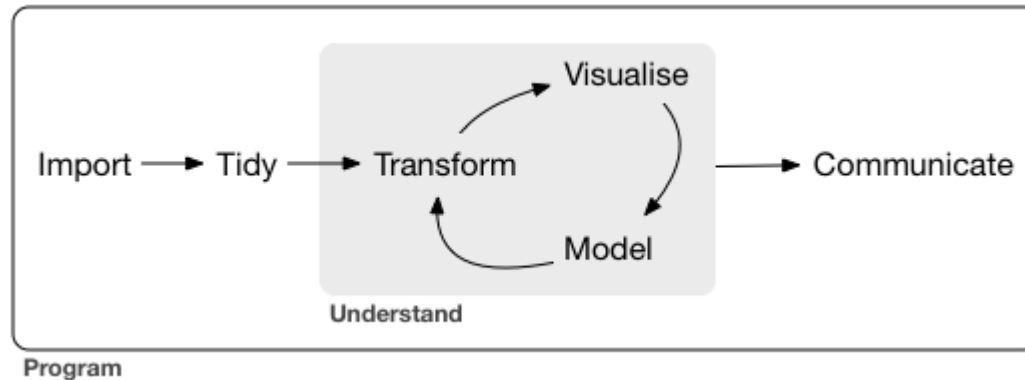
Exploring data

Into the tidyverse

June 6th, 2023

Data Science workflow

According to **Hadley Wickham** in **R for Data Science**:



First two weeks: data wrangling and visualization

Aspects of data **wrangling**:

- **import**: reading in data (e.g. `read_csv()`)
- **tidy**: rows = observations, columns = variables (i.e. **tabular** data)
- **transform**: filter observations, create new variables, summarize, etc.

What is Exploratory Data Analysis (EDA)?

(broadly speaking) EDA = questions about data + wrangling + visualization

R for Data Science: *"EDA is a state of mind"*, an iterative cycle:

- generate questions
- answer via transformations and visualizations

Example of questions?

- What type of **variation** do the variables display?
- What type of **relationships** exist between variables?

EDA is **NOT** a replacement for statistical inference and learning

EDA is an **important** and **necessary** step to build intuition

Now for an example...

Exploring MLB batting statistics

Import `Batting` table of historical MLB statistics from the `Lahman` package, explore using the `tidyverse`

```
library(tidyverse) # Load the tidyverse suite of packages
library(Lahman) # Load the Lahman package to access its datasets
Batting <- as_tibble(Batting) # Initialize the Batting dataset
```

Basic info about the `Batting` dataset:

```
dim(Batting) # displays same info as c(nrow(Batting), ncol(Batting))
```

```
## [1] 112184      22
```

```
class(Batting)
```

```
## [1] "tbl_df"      "tbl"        "data.frame"
```

`tbl` (pronounced `tibble`) is the `tidyverse` way of storing tabular data, like a spreadsheet or `data.frame`

Always look at your data: view the first 6 (by default) rows with `head()`

```
head(Batting) # Try just typing Batting into your console, what happens?
```

```
## # A tibble: 6 × 22
##   playerID yearID stint teamID lgID      G    AB    R    H   X2B   X3B
##   <chr>      <int> <int> <fct>  <fct> <int> <int> <int> <int> <int> <int> <
## 1 abercda01  1871     1  TRO    NA      1     4     0     0     0     0
## 2 addybo01   1871     1  RC1    NA     25    118    30    32     6     0
## 3 allisar01  1871     1  CL1    NA     29    137    28    40     4     5
## 4 allisdo01  1871     1  WS3    NA     27    133    28    44    10     2
## 5 ansonca01  1871     1  RC1    NA     25    120    29    39    11     3
## 6 armstbo01  1871     1  FW1    NA     12     49     9    11     2     1
## # ... with 10 more variables: RBI <int>, SB <int>, CS <int>, BB <int>, SO <in
## #   IBB <int>, HBP <int>, SH <int>, SF <int>, GIDP <int>
```

Is our **Batting** dataset **tidy**?

- Each row = a player's season stint with a team (i.e. players can play for multiple teams in year)
- Each column = different measurement or recording about the player-team-season observation (can print out column names directly with `colnames(Batting)` or `names(Batting)`)

Can we explore how baseball has changed over time with Batting?

Let the data wrangling begin...

Summarize *continuous* (e.g. `yearID`, `AB`) and *categorical* (e.g. `teamID`, `lgID`) variables in different ways

Compute **summary statistics** for *continuous* variables with the `summary()` function:

```
summary(Batting$yearID)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1871   1938   1978   1969   2003   2022
```

Compute **counts** of *categorical* variables with `table()` function:

```
table("Leagues" = Batting$lgID) # be careful it ignores NA values!
```

```
## Leagues
##      AA      AL      FL      NA      NL      PL      UA
##  1893 51799   472   737 56800   149   334
```

How do we remove the other leagues?

dplyr is a package within the **tidyverse** with functions for data wrangling

"Grammar of data manipulation": **dplyr** functions are **verbs**, datasets are **nouns**

- **We can `filter()` our dataset to choose observations meeting conditions**

```
mlb_batting <- filter(Batting, lgID %in% c("AL", "NL"))  
nrow(Batting) - nrow(mlb_batting) # Difference in rows
```

```
## [1] 3585
```

- **We can `select()` variables of interest**

```
sel_batting <- select(Batting, yearID, lgID, G, AB, R, H, HR, BB, SO)  
head(sel_batting, n = 3)
```

```
## # A tibble: 3 × 9  
##   yearID lgID      G    AB    R    H    HR    BB    SO  
##   <int> <fct> <int> <int> <int> <int> <int> <int> <int>  
## 1  1871 NA      1     4     0     0     0     0     0  
## 2  1871 NA     25    118    30    32     0     4     0  
## 3  1871 NA     29    137    28    40     0     2     5
```

- **We can arrange() our dataset to sort observations by variables**

```
hr_batting <- arrange(Batting, desc(HR)) # use desc() for descending order
head(hr_batting, n = 3)
```

```
## # A tibble: 3 × 22
##   playerID  yearID stint teamID lgID      G    AB    R    H   X2B   X3B
##   <chr>      <int> <int> <fct>  <fct> <int> <int> <int> <int> <int> <int> <int>
## 1 bondsba01  2001     1 SFN    NL    153  476  129  156   32    2
## 2 mcgwima01  1998     1 SLN    NL    155  509  130  152   21    0
## 3 sosasa01  1998     1 CHN    NL    159  643  134  198   20    0
## # ... with 10 more variables: RBI <int>, SB <int>, CS <int>, BB <int>, SO <int>,
## #   IBB <int>, HBP <int>, SH <int>, SF <int>, GIDP <int>
```

- **We can summarize() our dataset to one row based on functions of variables**

```
summarize(Batting, max(stint), median(AB))
```

```
## # A tibble: 1 × 2
##   `max(stint)` `median(AB)`
##   <int>      <dbl>
## 1           5         45
```


- **We can mutate() our dataset to create new variables** (mutate is a weird name...)

```
new_batting <- mutate(Batting, batting_avg = H / AB, so_to_bb = SO / BB)
head(new_batting, n = 1)
```

```
## # A tibble: 1 × 24
##   playerID yearID stint teamID lgID      G    AB    R    H   X2B   X3B
##   <chr>      <int> <int> <fct>  <fct> <int> <int> <int> <int> <int> <int> <
## 1 abercda01  1871     1 TRO    NA      1     4     0     0     0     0
## # ... with 12 more variables: RBI <int>, SB <int>, CS <int>, BB <int>, SO <in
## #   IBB <int>, HBP <int>, SH <int>, SF <int>, GIDP <int>, batting_avg <dbl>
## #   so_to_bb <dbl>
```

How do we perform several of these actions?

```
head(arrange(select(mutate(Batting, BA = H / AB), playerID, BA), desc(BA
```

```
## # A tibble: 1 × 2
##   playerID    BA
##   <chr>      <dbl>
## 1 snowch01     1
```

That's awfully annoying to do, and also difficult to read...

Enter the pipeline

The `%>%` (*pipe*) operator is used in the **tidyverse** (from **magrittr**) to chain commands together

`%>%` directs the **data analysis pipeline**: output of one function pipes into input of the next function

```
Batting %>%  
  filter(lgID %in% c("AL", "NL"),  
         AB > 300) %>%  
  mutate(batting_avg = H / AB) %>%  
  arrange(desc(batting_avg)) %>%  
  select(playerID, yearID, batting_avg) %>%  
  head(n = 5)
```

```
## # A tibble: 5 × 3  
##   playerID  yearID batting_avg  
##   <chr>      <int>      <dbl>  
## 1 duffyhu01   1894      0.440  
## 2 barnero01   1876      0.429  
## 3 lajoina01   1901      0.426  
## 4 keelewi01   1897      0.424  
## 5 hornsro01   1924      0.424
```

More pipeline actions!

Instead of `head()`, **we can `slice()` our dataset to choose the observations based on the position**

```
Batting %>%  
  filter(lgID %in% c("AL", "NL"),  
         AB > 300) %>%  
  mutate(so_to_bb = SO / BB) %>%  
  arrange(so_to_bb) %>%  
  select(playerID, yearID, so_to_bb) %>%  
  slice(c(1, 2, 10, 100))
```

```
## # A tibble: 4 × 3  
##   playerID  yearID so_to_bb  
##   <chr>      <int>    <dbl>  
## 1 roweja01    1882      0  
## 2 seweljo01   1932    0.0536  
## 3 holloch01   1922    0.0862  
## 4 collied01   1918    0.178
```

Grouped operations

We `group_by()` to split our dataset into groups based on a variable's values

```
Batting %>%  
  filter(lgID %in% c("AL", "NL")) %>%  
  group_by(yearID) %>%  
  summarize(hr = sum(HR), so = sum(SO), bb = sum(BB)) %>%  
  arrange(desc(hr)) %>%  
  slice(1:5)
```

```
## # A tibble: 5 × 4  
##   yearID    hr    so    bb  
##   <int> <int> <int> <int>  
## 1  2019  6776 42823 15895  
## 2  2017  6105 40104 15829  
## 3  2021  5944 42145 15794  
## 4  2000  5693 31356 18237  
## 5  2016  5610 38982 15088
```

`group_by()` is only useful in a pipeline (e.g. with `summarize()`), and pay attention to its behavior

`ungroup()` can solve your problems afterwards

Putting it all together...

We'll create a **tidy** dataset where each row = a year with the following variables:

- total HRs (homeruns), SOs (strikeouts), and BBs (walks)
- year's BA = total H / total AB
- only want AL and NL leagues

```
year_batting_summary <- Batting %>%  
  filter(lgID %in% c("AL", "NL")) %>%  
  group_by(yearID) %>%  
  summarize(total_hits = sum(H, na.rm = TRUE),  
            total_hrs = sum(HR, na.rm = TRUE),  
            total_sos = sum(SO, na.rm = TRUE),  
            total_walks = sum(BB, na.rm = TRUE),  
            total_atbats = sum(AB, na.rm = TRUE)) %>%  
  mutate(batting_avg = total_hits / total_atbats)  
head(year_batting_summary, n = 2)
```

```
## # A tibble: 2 × 7
```

##	yearID	total_hits	total_hrs	total_sos	total_walks	total_atbats	batting_av
##	<int>	<int>	<int>	<int>	<int>	<int>	<dbl>
## 1	1876	5338	40	589	336	20121	0.26
## 2	1877	3705	24	726	345	13667	0.27

Top three years with the most HRs?

```
year_batting_summary %>%  
  arrange(desc(total_hrs)) %>%  
  slice(1:3)
```

```
## # A tibble: 3 × 7
```

```
##   yearID total_hits total_hrs total_sos total_walks total_atbats batting_av  
##   <int>     <int>     <int>     <int>     <int>       <int>     <dbl>  
## 1  2019     42039     6776     42823     15895     166651     0.25  
## 2  2017     42215     6105     40104     15829     165567     0.25  
## 3  2021     39484     5944     42145     15794     161941     0.24
```

Top three years with highest batting average?

```
year_batting_summary %>%  
  arrange(desc(batting_avg)) %>%  
  slice(1:3)
```

```
## # A tibble: 3 × 7
```

```
##   yearID total_hits total_hrs total_sos total_walks total_atbats batting_av  
##   <int>     <int>     <int>     <int>     <int>       <int>     <dbl>  
## 1  1894     17809      629      3333      5870      57577     0.30  
## 2  1895     16827      488      3621      5120      56788     0.29  
## 3  1930     25597     1565      7934      7654      86571     0.29
```

Best and worst strikeout to walk ratios?

```
year_batting_summary %>%  
  mutate(so_to_bb = total_sos / total_walks) %>%  
  arrange(so_to_bb) %>%  
  slice(c(1, n()))
```

```
## # A tibble: 2 × 8  
##   yearID total_hits total_hrs total_sos total_walks total_atbats batti...1 so  
##   <int>      <int>      <int>      <int>      <int>      <int>      <dbl>  
## 1   1893      15913       460      3341       6143      56898      0.280  
## 2   1879       6171        58      1843        508      24155      0.255  
## # ... with abbreviated variable names 1batting_avg, 2so_to_bb
```

We can make better looking tables... **rename()** variables in our dataset

```
year_batting_summary %>%  
  select(yearID, batting_avg) %>%  
  rename(Year = yearID, `Batting AVG` = batting_avg) %>%  
  slice(c(1, n()))
```

```
## # A tibble: 2 × 2  
##   Year `Batting AVG`  
##   <int>      <dbl>  
## 1  1876      0.265  
## 2  2022      0.243
```

Grammar of tables preview

We can go one step further - **and use the new gt package** to create a nice-looking table for presentation

```
library(gt)
year_batting_summary %>%
  select(yearID, batting_avg) %>%
  rename(Year = yearID,
         `Batting AVG` = batting_avg) %>%
  arrange(desc(`Batting AVG`)) %>%
  slice(c(1:3, (n()-2):n())) %>%
  gt() %>%
  tab_header(
    title = "Best / worst MLB Seasons",
    subtitle = "Top / bottom three are presented"
  )
```

```
##
## Attaching package: 'gt'
##
## The following object is masked from 'base':
##
##     google_font
```

Best / worst MLB Seasons by AVG

Top / bottom three are presented

Year	Batting AVG
1894	0.3093075
1895	0.2963126
1930	0.2956764
1908	0.2389593
1888	0.2387601
1885	0.2386801

Data visualization

"The simple graph has brought more information to the data analyst's mind than any other device." — Tukey

- **TOMORROW**: the **grammar of graphics**
- Use `ggplot2` to visually explore our data
- More intuitive than base **R** plotting!
- Will walkthrough different types of visualizations for 1D, 2D, continuous, categorical, facetting, etc.
- **tidyverse** verbs and `%>%` leads to natural pipeline for EDA

