

Supervised Learning

Linear regression

June 21st, 2021

Simple linear regression

We assume a **linear relationship** for $Y = f(X)$:

$$Y_i = \beta_0 + \beta_1 X_i + \epsilon_i, \quad \text{for } i = 1, 2, \dots, n$$

- Y_i is the i th value for the **response** variable
- X_i is the i th value for the **predictor** variable
- β_0 is an *unknown*, constant **intercept**: average value for Y if $X = 0$
- β_1 is an *unknown*, constant **slope**: increase in average value for Y for each one-unit increase in X
- ϵ_i is the **random** noise: assume **independent, identically distributed** (*iid*) from Normal distribution

$$\epsilon_i \stackrel{iid}{\sim} N(0, \sigma^2) \quad \text{with constant variance } \sigma^2$$

Simple linear regression estimation

We are estimating the **conditional expectation (mean)** for Y :

$$\mathbb{E}[Y_i|X_i] = \beta_0 + \beta_1 X_i$$

- average value for Y given the value for X

How do we estimate the **best fitting** line?

Ordinary least squares (OLS) - by minimizing the **residual sum of squares (RSS)**

$$RSS(\beta_0, \beta_1) = \sum_{i=1}^n [Y_i - (\beta_0 + \beta_1 X_i)]^2 = \sum_{i=1}^n (Y_i - \beta_0 - \beta_1 X_i)^2$$

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sum_{i=1}^n (X_i - \bar{X})^2} \quad \text{and} \quad \hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{X}$$

where $\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$ and $\bar{Y} = \frac{1}{n} \sum_{i=1}^n Y_i$

Connection to covariance and correlation

Covariance describes the **joint variability of two variables**

$$\text{Cov}(X, Y) = \sigma_{X,Y} = \mathbb{E}[(X - \mathbb{E}[X])(Y - \mathbb{E}[Y])]$$

We compute the **sample covariance** (use $n - 1$ since we are using the means and want **unbiased estimates**)

$$\hat{\sigma}_{X,Y} = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})$$

Correlation is a *normalized* form of covariance, ranges from -1 to 1

$$\rho_{X,Y} = \frac{\text{Cov}(X, Y)}{\sigma_X \cdot \sigma_Y}$$

Sample correlation uses the sample covariance and standard deviations, e.g. $s_X^2 = \frac{1}{n-1} \sum_i (X_i - \bar{X})^2$

$$r_{X,Y} = \frac{\sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y})}{\sqrt{\sum_{i=1}^n (X_i - \bar{X})^2 \sum_{i=1}^n (Y_i - \bar{Y})^2}}$$

Connection to covariance and correlation

So we have the following:

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (X_i - \bar{X}) (Y_i - \bar{Y})}{\sum_{i=1}^n (X_i - \bar{X})^2} \quad \text{compared to} \quad r_{X,Y} = \frac{\sum_{i=1}^n (X_i - \bar{X}) (Y_i - \bar{Y})}{\sqrt{\sum_{i=1}^n (X_i - \bar{X})^2 \sum_{i=1}^n (Y_i - \bar{Y})^2}}$$

\Rightarrow Can rewrite $\hat{\beta}_1$ as:

$$\hat{\beta}_1 = r_{X,Y} \cdot \frac{s_Y}{s_X}$$

\Rightarrow Can rewrite $r_{X,Y}$ as:

$$r_{X,Y} = \hat{\beta}_1 \cdot \frac{s_X}{s_Y}$$

Can think of $\hat{\beta}_1$ weighting the ratio of variance between X and Y ...

Example data: NFL teams summary

Created dataset using `nflfastR` summarizing NFL team performances from 1999 to 2020

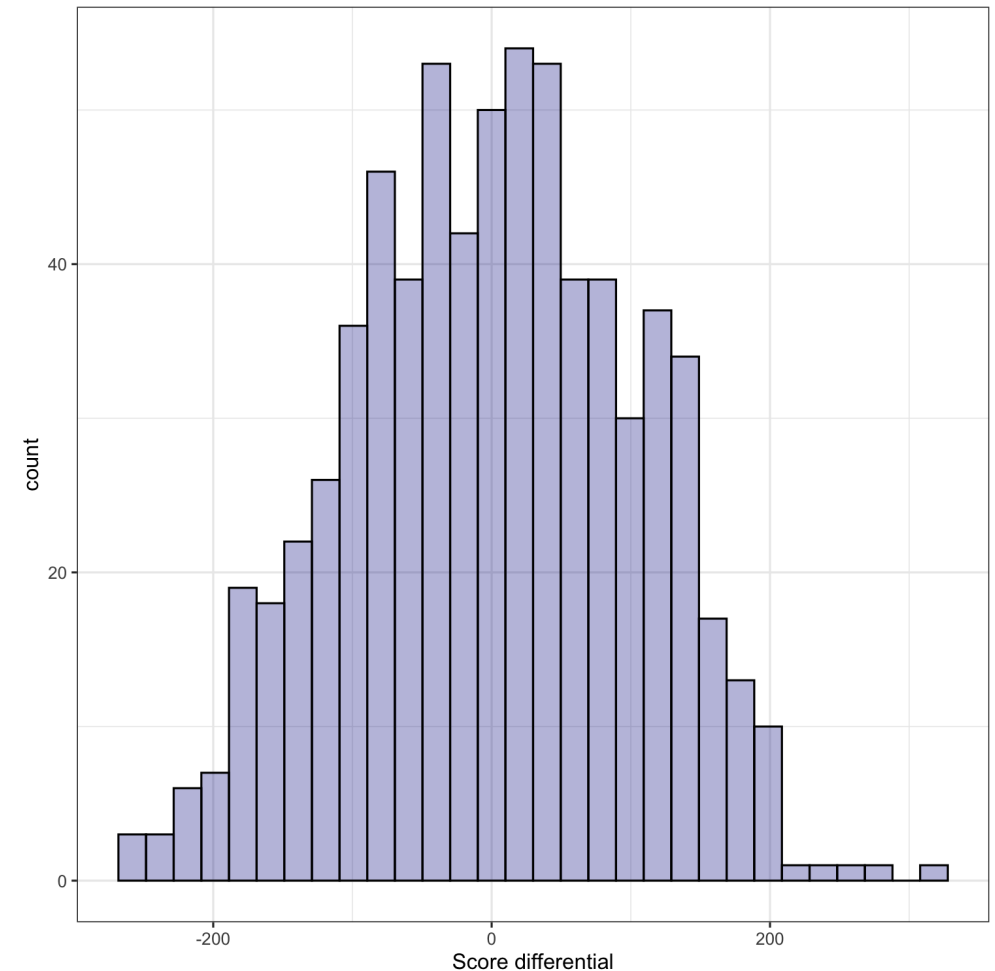
```
library(tidyverse)
nfl_teams_data <- read_csv("http://www.stat.cmu.edu/cmsac/sure/2021/materials/data/regression_pro
nfl_teams_data
```

```
## # A tibble: 701 × 55
##   season team  offens...1 offen...2 offen...3 offen...4 offen...5 offen...6 offen...7 offen...8
##   <dbl> <chr>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
## 1  1999 ARI      0.477    2796    1209    4.67    3.15      0      NaN      0
## 2  1999 ATL      0.504    3317    1176    6.08    3.20      0      NaN     11
## 3  1999 BAL      0.452    2805    1663    5.07    4.13      0      NaN      0
## 4  1999 BUF      0.540    3275    2038    6.17    4.13      0      NaN    161
## 5  1999 CAR      0.552    4144    1484    6.68    4.29      0      NaN     89
## 6  1999 CHI      0.561    4090    1359    5.75    3.55      0      NaN    508
## 7  1999 CIN      0.498    3178    1971    5.37    4.63      0      NaN      0
## 8  1999 CLE      0.489    2574    1140    4.71    3.67      0      NaN     35
## 9  1999 DAL      0.560    3083    2054    5.95    4.29      0      NaN      0
## 10 1999 DEN      0.546    3378    1852    5.85    4.05      0      NaN      9
## # ... with 691 more rows, 45 more variables: offense_ave_yac <dbl>,
## #   offense_n_plays_pass <dbl>, offense_n_plays_run <dbl>,
## #   offense_n_interceptions <dbl>, offense_n_fumbles_lost_pass <dbl>,
## #   offense_n_fumbles_lost_run <dbl>, offense_total_epa_pass <dbl>.
```

Modeling NFL score differential

Interested in modeling a team's **score differential**

```
nfl_teams_data <- nfl_teams_data %>%  
  mutate(score_diff =  
    points_scored - points_allowed)  
nfl_teams_data %>%  
  ggplot(aes(x = score_diff)) +  
    geom_histogram(color = "black",  
                  fill = "darkblue",  
                  alpha = 0.3) +  
  theme_bw() +  
  labs(x = "Score differential")
```




Relationship between score differential and passing performance

- offense_ave_epa_pass: team's average **expected points added (EPA)** per pass attempt

```
nfl_plot <- nfl_teams_data %>%  
  ggplot(aes(x = offense_ave_epa_pass,  
             y = score_diff)) +  
  geom_point(alpha = 0.5) +  
  theme_bw() +  
  labs(x = "EPA gained per pass attempt",  
       y = "Score differential")  
nfl_plot
```

We fit linear regression models using `lm()`,
formula is input as: response ~ predictor

```
init_lm <- lm(score_diff ~ offense_ave_epa_pa  
              data = nfl_teams_data)
```



View the model summary()

```
summary(init_lm)
```

```
##
## Call:
## lm(formula = score_diff ~ offense_ave_epa_pass, data = nfl_teams_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -182.611  -46.600   -1.326   44.123  233.640
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)      -4.890       2.561  -1.909   0.0566 .
## offense_ave_epa_pass  566.352     19.226  29.458  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 67.67 on 699 degrees of freedom
## Multiple R-squared:  0.5539,    Adjusted R-squared:  0.5532
## F-statistic: 867.8 on 1 and 699 DF,  p-value: < 2.2e-16
```

Inference with OLS

Reports the intercept and coefficient estimates: $\hat{\beta}_0 \approx -4.89$, $\hat{\beta}_1 \approx 566.352$

Estimates of uncertainty for β s via **standard errors**: $\widehat{SE}(\hat{\beta}_0) \approx 2.561$, $\widehat{SE}(\hat{\beta}_1) \approx 19.226$

t -statistics are coefficients Estimates / Std. Error, i.e., number of standard deviations from 0

- p -values (i.e., $\Pr(>|t|)$): estimated probability observing value as extreme as $|t \text{ value}|$ **given the null hypothesis** $\beta = 0$
- $p\text{-value} < \text{conventional threshold of } \alpha = 0.05$, **sufficient evidence to reject the null hypothesis that the coefficient is zero**,
- Typically $|t \text{ values}| > 2$ indicate **significant** relationship at $\alpha = 0.05$
- i.e., there is a **significant** association between offense_ave_epa_pass and score_diff

Be careful!

Caveats to keep in mind regarding p-values:

If the true value of a coefficient $\beta = 0$, then the p-value is sampled from a **Uniform(0,1) distribution**

- i.e. it is just as likely to have value 0.45 as 0.16 or 0.84 or 0.9999 or 0.00001...

⇒ Hence why we typically only reject for low α values like 0.05

- Controlling the Type 1 error rate at $\alpha = 0.05$, i.e., the probability of a **false positive** mistake
- 5% chance that you'll conclude there's a significant association between x and y **even when there is none**

Remember what a standard error is? $SE = \frac{\sigma}{\sqrt{n}}$

- ⇒ As n gets large **standard error goes to zero**, and *all* predictors are eventually deemed significant
- While the p-values might be informative, we will explore other approaches to determine which subset of predictors to include (e.g., holdout performance)

Back to the model summary: Multiple R-squared

Back to the connection between the coefficient and correlation:

$$r_{X,Y} = \hat{\beta}_1 \cdot \frac{s_X}{s_Y} \quad \Rightarrow \quad r_{X,Y}^2 = \hat{\beta}_1^2 \cdot \frac{s_X^2}{s_Y^2}$$

Compute the correlation with `cor()`:

```
with(nfl_teams_data, cor(offense_ave_epa_pass, score_diff))
```

```
## [1] 0.7442135
```

The squared cor matches the reported Multiple R-squared

```
with(nfl_teams_data, cor(offense_ave_epa_pass, score_diff))^2
```

```
## [1] 0.5538537
```

Back to the model summary: Multiple R-squared

Back to the connection between the coefficient and correlation:

$$r_{X,Y} = \hat{\beta}_1 \cdot \frac{s_X}{s_Y} \quad \Rightarrow \quad r_{X,Y}^2 = \hat{\beta}_1^2 \cdot \frac{s_X^2}{s_Y^2}$$

r^2 (or also R^2) estimates the **proportion of the variance** of Y explained by X

- More generally: variance of model predictions / variance of Y

```
var(predict(init_lm)) / var(nfl_teams_data$score_diff)
```

```
## [1] 0.5538537
```

Generating predictions

We can use the `predict()` function to either get the fitted values of the regression:

```
train_preds <- predict(init_lm)
head(train_preds)
```

```
##           1           2           3           4           5           6
## -120.21628  -33.58829 -104.31202   21.15045   57.18906  -23.34489
```

Which is equivalent to using:


```
head(init_lm$fitted.values)
```

```
##           1           2           3           4           5           6
## -120.21628  -33.58829 -104.31202   21.15045   57.18906  -23.34489
```

Predictions for new data

Or we can provide it newdata which **must contain the explanatory variables**:

```
steelers_data <- nfl_teams_data %>%  
  filter(team == "PIT", season == 2020)  
  
new_steelers_data <- steelers_data %>%  
  dplyr::select(team, offense_ave_epa_pass) %  
  slice(rep(1, 3)) %>%  
  mutate(adj_factor = c(0.5, 1, 2),  
         offense_ave_epa_pass = offense_ave_e  
new_steelers_data$pred_score_diff <-  
  predict(init_lm, newdata = new_steelers_dat  
nfl_plot +  
  geom_point(data = new_steelers_data,  
            aes(x = offense_ave_epa_pass,  
                y = pred_score_diff),  
            color = "gold", size = 5)
```




Plot observed values against predictions

Useful diagnostic (for **any type of model**, not just linear regression!)

```
nfl_teams_data %>%  
  mutate(pred_vals = predict(init_lm)) %>%  
  ggplot(aes(x = pred_vals,  
             y = score_diff)) +  
  geom_point(alpha = 0.5) +  
  geom_abline(slope = 1, intercept = 0,  
             linetype = "dashed",  
             color = "red",  
             size = 2) +  
  theme_bw()
```

- "Perfect" model will follow **diagonal**




Plot observed values against predictions

Can augment the data with model output using the **broom** package

```
nfl_teams_data <-  
  broom::augment(init_lm, nfl_teams_data)  
nfl_teams_data %>%  
  ggplot(aes(x = .fitted,  
             y = score_diff)) +  
  geom_point(alpha = 0.5) +  
  geom_abline(slope = 1, intercept = 0,  
             linetype = "dashed",  
             color = "red",  
             size = 2) +  
  theme_bw()
```

- Adds various columns from model fit we can use in plotting for model diagnostics




Plot residuals against predicted values

- Residuals = observed - predicted
- Conditional on the predicted values, the **residuals should have a mean of zero**

```
nfl_teams_data %>%  
  ggplot(aes(x = .fitted,  
             y = .resid)) +  
  geom_point(alpha = 0.5) +  
  geom_hline(yintercept = 0,  
            linetype = "dashed",  
            color = "red",  
            size = 2) +  
  # To plot the residual mean  
  geom_smooth(se = FALSE) +  
  theme_bw()
```

- Residuals **should NOT display any pattern**



Multiple regression

We can include as many variables as we want (assuming $n > p$!)

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p + \epsilon$$

OLS estimates in matrix notation (\mathbf{X} is a $n \times p$ matrix):

$$\hat{\boldsymbol{\beta}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$$

Can just add more variables to the formula in R

```
multiple_lm <- lm(score_diff ~ offense_ave_epa_pass + defense_ave_epa_pass,  
                  data = nfl_teams_data)
```

- Use the Adjusted R-squared when including multiple variables $= 1 - \frac{(1-R^2)(n-1)}{(n-p-1)}$
 - Adjusts for the number of variables in the model p
 - Adding more variables **will always increase** Multiple R-squared

What about the Normal distribution assumption???

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p + \epsilon$$

- ϵ_i is the **random** noise: assume **independent, identically distributed (iid)** from Normal distribution

$$\epsilon_i \stackrel{iid}{\sim} N(0, \sigma^2) \quad \text{with constant variance } \sigma^2$$

OLS doesn't care about this assumption, it's just estimating coefficients!

In order to perform inference, **we need to impose additional assumptions**

By assuming $\epsilon_i \stackrel{iid}{\sim} N(0, \sigma^2)$, what we really mean is:

$$Y \stackrel{iid}{\sim} N(\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p, \sigma^2)$$


So we're estimating the mean μ of this conditional distribution, but what about σ^2 ?

Unbiased estimate $\hat{\sigma}^2 = \frac{RSS}{n-(p+1)}$, its square root is the Residual standard error

- **Degrees of freedom:** $n - (p + 1)$, data supplies us with n "degrees of freedom" and we used up $p + 1$

Check the assumptions about normality with `ggfortify`

```
library(ggfortify)
autoplot(multiple_lm, ncol = 4) + theme_bw()
```



- Standardized residuals = $\text{residuals} / \text{sd}(\text{residuals})$ (see also `.std.resid` from `augment`)