

Advanced topics

More fun with classification

July 22nd, 2021

NFL completion probability

Binary outcome model: $Y \in \{\text{Incomplete (0), Complete (1)}\}$

```
library(tidyverse)
nfl_passing_plays <-
  read_csv("http://www.stat.cmu.edu/cmsac/sure/2021/materials/data/eda_projects/nfl_passing_plays.csv")
  # Only keep rows with passer and receiver information known:
  filter(!is.na(passer_player_id), !is.na(receiver_player_id),
         !is.na(epa), !is.na(air_yards), !is.na(pass_location)) %>%
  # Combine passer and receiver unique IDs:
  mutate(passer_name_id = paste0(passer_player_name, ":", passer_player_id),
         receiver_name_id = paste0(receiver_player_name, ":", receiver_player_id))
```

Create train and test folds based on games:

```
set.seed(1985)
game_fold_table <- tibble(game_id = unique(nfl_passing_plays$game_id)) %>%
  mutate(game_fold = sample(rep(1:5, length.out = n()), n()))
nfl_passing_plays <- nfl_passing_plays %>% dplyr::left_join(game_fold_table, by = "game_id")
```

Logistic regression review

Generate data of test predictions with particular model:

```
logit_cv_preds <-  
  map_dfr(unique(nfl_passing_plays$game_fold),  
    function(test_fold) {  
      # Separate test and training data:  
      test_data <- nfl_passing_plays %>%  
        filter(game_fold == test_fold)  
      train_data <- nfl_passing_plays %>%  
        filter(game_fold != test_fold)  
  
      # Train model:  
      logit_model <- glm(complete_pass ~ yardline_100 + shotgun + air_yards + pass_location,  
        data = train_data, family = "binomial")  
  
      # Return tibble of holdout results:  
      tibble(test_pred_probs = predict(logit_model, newdata = test_data,  
        type = "response"),  
        test_actual = test_data$complete_pass,  
        game_fold = test_fold)  
    })
```

Holdout performance by fold

```
logit_cv_preds %>%  
  mutate(test_pred = ifelse(test_pred_probs < .5, 0, 1)) %>%  
  group_by(game_fold) %>%  
  summarize(mcr = mean(test_pred != test_actual))
```

```
## # A tibble: 5 × 2  
##   game_fold  mcr  
##   <int> <dbl>  
## 1      1 0.297  
## 2      2 0.298  
## 3      3 0.293  
## 4      4 0.277  
## 5      5 0.296
```

Let's think more carefully about what's going on here...

Evaluating the prediction threshold

We can really write our classification as a function of some cutoff c :

$$\hat{Y} = \hat{C}(x) = \begin{cases} 1 & \hat{p}(x) > c \\ 0 & \hat{p}(x) \leq c \end{cases}$$

Given the classifications, we can form a confusion matrix:

	predicted events	predicted non-events		predicted events	predicted non-events
actual events	correctly forecasted events	missed events	actual events	True Positive	False Negative
actual non-events	missed non-events	correctly forecasted non-events	actual non-events	False Positive	True Negative



We want to **maximize** all of the following (positive means 1, negative means 0):

- **Accuracy:** How often is the classifier correct? $\frac{TP+TN}{total}$
- **Precision:** How often is it right for predicted positives? $\frac{TP}{TP+FP}$
- **Sensitivity**, aka true positive rate (TPR) or power: How often does it detect positives? $\frac{TP}{TP+FN}$
- **Specificity**, aka true negative rate (TNR), or 1 - false positive rate (FPR): How often does it detect negatives? $\frac{TN}{TN+FP}$

So how do we handle this?


We want to balance with high power and low false positive rate



Receiver Operating Characteristic (ROC) curve

Check all possible values for the cutoff c , plot the power against false positive rate


Want to maximize the **area under the curve (AUC)**



plotROC and holdout AUC

- d stands for disease status (the outcome)
- m stands for marker (the prediction)

```
library(plotROC)
logit_cv_preds %>%
  ggplot() +
  geom_roc(aes(d = test_actual,
               m = test_pred_probs),
            labelround = 4) +
  style_roc() +
  geom_abline(slope = 1, intercept = 0, linet
  labs(color = "Test fold")
with(logit_cv_preds,
      MLmetrics::AUC(test_pred_probs, test_act
```




```
## [1] 0.6864571
```

plotROC and holdout AUC by test fold

```
logit_cv_preds %>%  
  ggplot() +  
  geom_roc(aes(d = test_actual,  
              m = test_pred_probs,  
              color = as.factor(game_fold)),  
           n.cuts = 0) +  
  style_roc() +  
  geom_abline(slope = 1, intercept = 0, linet  
  ggthemes::scale_color_colorblind() +  
  labs(color = "Test fold") +  
  theme(legend.position = "bottom")  
logit_cv_preds %>% group_by(game_fold) %>%  
  summarize(auc = MLmetrics::AUC(test_pred_pr
```

There is definitely room for improvement...



```
## # A tibble: 5 × 2  
##   game_fold auc  
##   <int> <dbl>  
## 1         1 0.690  
## 2         2 0.679  
## 3         3 0.678  
## 4         4 0.705  
## 5         5 0.681
```

Add in varying intercepts with glmer?

```
library(lme4)
glmer_cv_preds <-
  map_dfr(unique(nfl_passing_plays$game_fold),
    function(test_fold) {

      # Separate test and training data - scale variables:
      test_data <- nfl_passing_plays %>% filter(game_fold == test_fold) %>%
        mutate(yardline_100 = scale(yardline_100), air_yards = scale(air_yards))

      train_data <- nfl_passing_plays %>% filter(game_fold != test_fold) %>%
        mutate(yardline_100 = scale(yardline_100), air_yards = scale(air_yards))


      # Train model:
      glmer_model <- glmer(complete_pass ~ yardline_100 + shotgun + air_yards +
        pass_location + qb_hit +
        (1|passer_name_id) + (1|receiver_name_id),
        data = train_data, family = "binomial")

      # Return tibble of holdout results:
      tibble(test_pred_probs = predict(glmer_model, newdata = test_data,
        type = "response",
        allow.new.levels = TRUE),
        test_actual = test_data$complete_pass,
        game_fold = test_fold)
```

Add in varying intercepts with `glmer`?

```
glmer_cv_preds %>%  
  ggplot() +  
  geom_roc(aes(d = test_actual,  
              m = test_pred_probs,  
              color = as.factor(game_fold)),  
           n.cuts = 0) +  
  style_roc() +  
  geom_abline(slope = 1, intercept = 0, linet  
  ggthemes::scale_color_colorblind() +  
  labs(color = "Test fold") +  
  theme(legend.position = "bottom")  
glmer_cv_preds %>% group_by(game_fold) %>%  
  summarize(auc = MLmetrics::AUC(test_pred_pr
```

Looks like player-level effects do not help!



```
## # A tibble: 5 × 2  
##   game_fold   auc  
##   <int> <dbl>  
## 1         1 0.692  
## 2         2 0.681  
## 3         3 0.682  
## 4         4 0.705  
## 5         5 0.684
```

Tree-based approach?

We need to first convert categorical variables into dummy indicators:

```
model_data <- nfl_passing_plays %>%  
  mutate(play_id = 1:n(),  
         complete_pass = as.factor(complete_pass)) %>%  
  dplyr::select(play_id, complete_pass, yardline_100, shotgun, air_yards, qb_hit,  
               game_fold, pass_location) %>%  
  mutate(pass_location_val = 1) %>%  
  pivot_wider(id_cols = play_id:game_fold,  
              names_from = pass_location, values_from = pass_location_val,  
              values_fill = 0) %>%  
  dplyr::select(-play_id)
```

Random forests using probability forest

For each tree compute class proportion in terminal node, then take average across all trees

```
library(ranger)
rf_prob_cv_preds <-
  map_dfr(unique(model_data$game_fold),
    function(test_fold) {

      # Separate test and training data - scale variables:
      test_data <- model_data %>% filter(game_fold == test_fold)

      train_data <- model_data %>% filter(game_fold != test_fold)

      rf_prob_model <-
        ranger(complete_pass ~ ., data = dplyr::select(train_data, -game_fold),
              probability = TRUE)


      # Return tibble of holdout results:
      tibble(test_pred_probs =
        as.numeric(predict(rf_prob_model, data = test_data,
                          type = "response")$predictions[,2]),
        test_actual = as.numeric(test_data$complete_pass) - 1,
        game_fold = test_fold)

    })
```

Random forests using probability forest

```
rf_prob_cv_preds %>%  
  ggplot() +  
  geom_roc(aes(d = test_actual,  
              m = test_pred_probs,  
              color = as.factor(game_fold)),  
           n.cuts = 0) +  
  style_roc() +  
  geom_abline(slope = 1, intercept = 0, linet  
  ggthemes::scale_color_colorblind() +  
  labs(color = "Test fold") +  
  theme(legend.position = "bottom")  
rf_prob_cv_preds %>% group_by(game_fold) %>%  
  summarize(auc = MLmetrics::AUC(test_pred_pr
```

Looks like just a modest improvement



```
## # A tibble: 5 × 2  
##   game_fold auc  
##   <int> <dbl>  
## 1         1 0.702  
## 2         2 0.703  
## 3         3 0.683  
## 4         4 0.713  
## 5         5 0.690
```

XGBoost!

```
library(xgboost)
xgb_cv_preds <-
  map_dfr(unique(model_data$game_fold),
    function(test_fold) {
      # Separate test and training data - scale variables:
      test_data <- model_data %>% filter(game_fold == test_fold)
      test_data_x <- as.matrix(dplyr::select(test_data, -complete_pass, -game_fold))
      train_data <- model_data %>% filter(game_fold != test_fold)
      train_data_x <- as.matrix(dplyr::select(train_data, -complete_pass, -game_fold))
      train_data_y <- as.numeric(train_data$complete_pass) - 1


      xgb_model <- xgboost(data = train_data_x, label = train_data_y,
        nrounds = 100, max_depth = 3, eta = 0.3,
        gamma = 0, colsample_bytree = 1, min_child_weight = 1,
        subsample = 1, nthread = 1,
        objective = 'binary:logistic', eval_metric = 'auc',
        verbose = 0)

      # Return tibble of holdout results:
      tibble(test_pred_probs =
        as.numeric(predict(xgb_model, newdata = test_data_x, type = "response")),
        test_actual = as.numeric(test_data$complete_pass) - 1,
        game_fold = test_fold)
```


XGBoost

```
xgb_cv_preds %>%  
  ggplot() +  
  geom_roc(aes(d = test_actual,  
              m = test_pred_probs,  
              color = as.factor(game_fold)),  
           n.cuts = 0) +  
  style_roc() +  
  geom_abline(slope = 1, intercept = 0, linet  
  ggthemes::scale_color_colorblind() +  
  labs(color = "Test fold") +  
  theme(legend.position = "bottom")  
xgb_cv_preds %>% group_by(game_fold) %>%  
  summarize(auc = MLmetrics::AUC(test_pred_pr
```

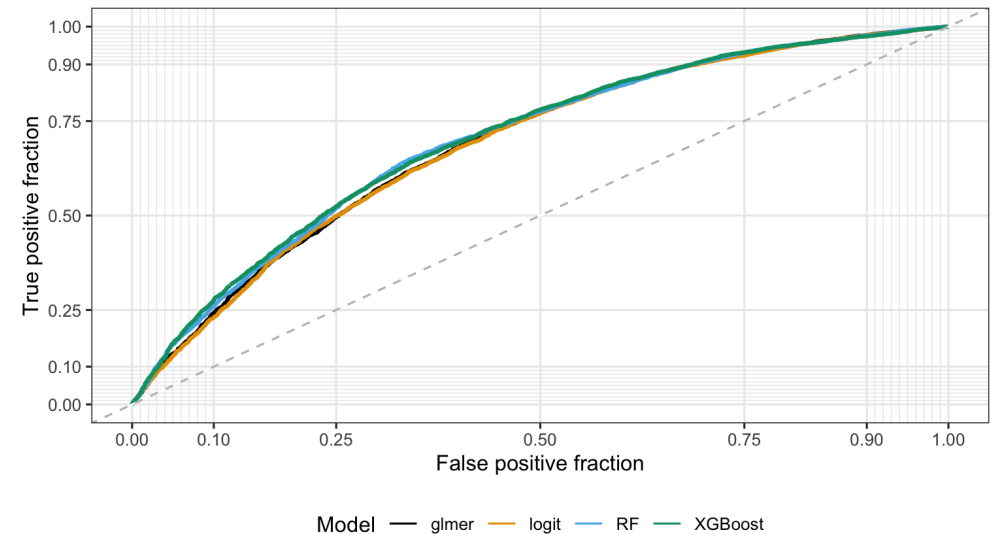
Should actually tune this more...



```
## # A tibble: 5 × 2  
##   game_fold auc  
##   <int> <dbl>  
## 1         1 0.704  
## 2         2 0.706  
## 3         3 0.683  
## 4         4 0.715  
## 5         5 0.692
```

All together now...

```
bind_rows(  
  mutate(logit_cv_preds, type = "logit"),  
  mutate(glmer_cv_preds, type = "glmer"),  
  mutate(rf_prob_cv_preds, type = "RF"),  
  mutate(xgb_cv_preds, type = "XGBoost")) %>%  
  ggplot() +  
  geom_roc(aes(d = test_actual,  
              m = test_pred_probs,  
              color = type),  
          n.cuts = 0) +  
  style_roc() +  
  geom_abline(slope = 1, intercept = 0, linet  
  ggthemes::scale_color_colorblind() +  
  labs(color = "Model") +  
  theme(legend.position = "bottom")
```



Pretty similar performance across all models...

Explaining predictions with SHAP-values

SHAP-values are based on **Shapley values** (an idea from game theory) and are used to measure the contributions from each feature in the model to the prediction for an individual observation

Shapley value ϕ_i^j for feature value j for observation i can be interpreted as:

- the value of feature j contributed ϕ_i^j to the prediction of observation i compared to the average prediction for the dataset
- linear regression coefficients function in the same way


Can use them in multiple ways:

- View total importance: $\frac{1}{n} \sum |\phi_i^j|$
- View distribution of ϕ_i^j for each feature
- Plot ϕ_i^j against feature value for partial dependence

SHAPforxgboost

Fit model on full data then **extract SHAP-values with SHAPforxgboost**

```
library(xgboost)
train_data_x <- as.matrix(dplyr::select(model_data, -complete_pass, -game_fold))
train_data_y <- as.numeric(model_data$complete_pass) - 1
xgb_model <- xgboost(data = train_data_x, label = train_data_y, nrounds = 100, max_depth = 3, eta = 0.1,
                     gamma = 0, colsample_bytree = 1, min_child_weight = 1, subsample = 1, nthread = 4,
                     objective = 'binary:logistic', eval_metric = 'auc', verbose = 0)
library(SHAPforxgboost)
shap_value_list <- shap.values(xgb_model, X_train = train_data_x)
shap.plot.summary.wrap1(xgb_model, X = train_data_x)
```



Multinomial classification with XGBoost

Use same NFL play-by-play dataset as before but get ready for XGBoost...

```
nfl_ep_model_data <- read_rds("http://www.stat.cmu.edu/cmsac/sure/2021/materials/data/model_pbp_c")
  mutate(Next_Score_Half = fct_relevel(Next_Score_Half,
                                       "No_Score", "Safety", "Field_Goal", "Touchdown",
                                       "Opp_Safety", "Opp_Field_Goal", "Opp_Touchdown"),
         next_score_label = as.numeric(Next_Score_Half) - 1)
model_variables <- c("half_seconds_remaining", "yardline_100", "down", "ydstogo")
```

XGBoost requires the multinomial categories to be numeric starting at 0

Leave-one-season-out cross-validation

```
library(xgboost)
xgb_loso_cv_preds <-
  map_dfr(unique(nfl_ep_model_data$season), function(x) {

    # Separate test and training data - scale variables:
    test_data <- nfl_ep_model_data %>% filter(season == x)
    test_data_x <- as.matrix(dplyr::select(test_data, model_variables))
    train_data <- nfl_ep_model_data %>% filter(season != x)
    train_data_x <- as.matrix(dplyr::select(train_data, model_variables))
    train_data_y <- train_data$next_score_label

    xgb_model <- xgboost(data = train_data_x, label = train_data_y, nrounds = 100, max_depth = 4,
                        eta = 0.3, gamma = 0, colsample_bytree = 1, min_child_weight = 1,
                        subsample = 1, nthread = 1, objective = 'multi:softprob', num_class = 7,
                        eval_metric = 'mlogloss', verbose = 0)

    xgb_preds <- matrix(predict(xgb_model, test_data_x), ncol = 7, byrow = TRUE) %>%
      as_tibble()
    colnames(xgb_preds) <- c("No_Score", "Safety", "Field_Goal", "Touchdown",
                           "Opp_Safety", "Opp_Field_Goal", "Opp_Touchdown")

    xgb_preds %>%
      mutate(Next_Score_Half = test_data$Next_Score_Half, season = x)
  })
```

Calibration results for each scoring event

```
ep_cv_loso_calibration_results <- xgb_loso_cv_preds %>%
  pivot_longer(No_Score:Opp_Touchdown,
               names_to = "next_score_type",
               values_to = "pred_prob") %>%
  mutate(bin_pred_prob = round(pred_prob / 0.05) * .05) %>%
  group_by(next_score_type, bin_pred_prob) %>%
  summarize(n_plays = n(),
            n_scoring_event = length(which(Next_Score_Half == next_score_type)),
            bin_actual_prob = n_scoring_event / n_plays,
            bin_se = sqrt((bin_actual_prob * (1 - bin_actual_prob)) / n_plays)) %>%
  ungroup() %>%
  mutate(bin_upper = pmin(bin_actual_prob + 2 * bin_se, 1),
         bin_lower = pmax(bin_actual_prob - 2 * bin_se, 0))
```

Calibration results for each scoring event

```
ep_cv_loso_calibration_results %>%
  mutate(next_score_type = fct_relevel(next_score_type, "Opp_Safety", "Opp_Field_Goal",
                                       "Opp_Touchdown", "No_Score", "Safety", "Field_Goal", "Touchdown"),
         next_score_type = fct_recode(next_score_type, "-Field Goal (-3)" = "Opp_Field_Goal", "-Safety (-3)" = "Opp_Safety",
                                       "Field Goal (3)" = "Field_Goal", "No Score (0)" = "No_Score",
                                       "Touchdown (7)" = "Touchdown", "Safety (2)" = "Safety")) %>%
  ggplot(aes(x = bin_pred_prob, y = bin_actual_prob)) +
  geom_abline(slope = 1, intercept = 0, color = "black", linetype = "dashed") +
  geom_smooth(se = FALSE) +
  geom_point(aes(size = n_plays)) +
  geom_errorbar(aes(ymin = bin_lower, ymax = bin_upper)) + #coord_equal() +
  scale_x_continuous(limits = c(0,1)) +
  scale_y_continuous(limits = c(0,1)) +
  labs(size = "Number of plays", x = "Estimated next score probability",
       y = "Observed next score probability") +
  theme_bw() +
  theme(strip.background = element_blank(),
        axis.text.x = element_text(angle = 90),
        legend.position = c(1, .05), legend.justification = c(1, 0)) +
  facet_wrap(~ next_score_type, ncol = 4)
```


Calibration results for each scoring event

